

```

#Importing the required libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn import metrics

data = pd.read_csv('uber.csv')

data

      Unnamed: 0           key  fare_amount \
0        24238194  2015-05-07 19:52:06.0000003    7.5
1        27835199  2009-07-17 20:04:56.0000002    7.7
2        44984355  2009-08-24 21:45:00.00000061   12.9
3        25894730  2009-06-26 08:22:21.0000001    5.3
4        17610152  2014-08-28 17:47:00.000000188   16.0
...
199995     ...  2012-10-28 10:49:00.00000053    3.0
199996     42598914  2014-03-14 01:09:00.0000008    7.5
199997     16382965  2009-06-29 00:42:00.00000078   30.9
199998     27804658  2015-05-20 14:56:25.0000004   14.5
199999     20259894  2010-05-15 04:08:00.00000076   14.1

      pickup_datetime  pickup_longitude  pickup_latitude \
0  2015-05-07 19:52:06 UTC          -73.999817       40.738354
1  2009-07-17 20:04:56 UTC          -73.994355       40.728225
2  2009-08-24 21:45:00 UTC          -74.005043       40.740770
3  2009-06-26 08:22:21 UTC          -73.976124       40.790844
4  2014-08-28 17:47:00 UTC          -73.925023       40.744085
...
199995     ...  2012-10-28 10:49:00 UTC          -73.987042       40.739367
199996     2014-03-14 01:09:00 UTC          -73.984722       40.736837
199997     2009-06-29 00:42:00 UTC          -73.986017       40.756487
199998     2015-05-20 14:56:25 UTC          -73.997124       40.725452
199999     2010-05-15 04:08:00 UTC          -73.984395       40.720077

      dropoff_longitude  dropoff_latitude  passenger_count
0          -73.999512       40.723217                  1
1          -73.994710       40.750325                  1
2          -73.962565       40.772647                  1
3          -73.965316       40.803349                  3
4          -73.973082       40.761247                  5
...
199995     ...  -73.986525       40.740297                  1
199996     -74.006672       40.739620                  1

```

```

199997      -73.858957      40.692588      2
199998      -73.983215      40.695415      1
199999      -73.985508      40.768793      1

```

[200000 rows x 9 columns]

data.head()

	Unnamed: 0	key	fare_amount	\
0	24238194	2015-05-07 19:52:06.0000003		7.5
1	27835199	2009-07-17 20:04:56.0000002		7.7
2	44984355	2009-08-24 21:45:00.00000061		12.9
3	25894730	2009-06-26 08:22:21.0000001		5.3
4	17610152	2014-08-28 17:47:00.000000188		16.0
		pickup_datetime	pickup_longitude	pickup_latitude
0	2015-05-07 19:52:06 UTC	-73.999817	40.738354	\
1	2009-07-17 20:04:56 UTC	-73.994355	40.728225	
2	2009-08-24 21:45:00 UTC	-74.005043	40.740770	
3	2009-06-26 08:22:21 UTC	-73.976124	40.790844	
4	2014-08-28 17:47:00 UTC	-73.925023	40.744085	
		dropoff_longitude	dropoff_latitude	passenger_count
0	-73.999512	40.723217	1	
1	-73.994710	40.750325	1	
2	-73.962565	40.772647	1	
3	-73.965316	40.803349	3	
4	-73.973082	40.761247	5	

data.tail()

	Unnamed: 0	key	fare_amount	\
199995	42598914	2012-10-28 10:49:00.00000053		3.0
199996	16382965	2014-03-14 01:09:00.0000008		7.5
199997	27804658	2009-06-29 00:42:00.00000078		30.9
199998	20259894	2015-05-20 14:56:25.0000004		14.5
199999	11951496	2010-05-15 04:08:00.00000076		14.1
		pickup_datetime	pickup_longitude	pickup_latitude
199995	2012-10-28 10:49:00 UTC	-73.987042	40.739367	\
199996	2014-03-14 01:09:00 UTC	-73.984722	40.736837	
199997	2009-06-29 00:42:00 UTC	-73.986017	40.756487	
199998	2015-05-20 14:56:25 UTC	-73.997124	40.725452	
199999	2010-05-15 04:08:00 UTC	-73.984395	40.720077	
		dropoff_longitude	dropoff_latitude	passenger_count
199995	-73.986525	40.740297	1	
199996	-74.006672	40.739620	1	
199997	-73.858957	40.692588	2	
199998	-73.983215	40.695415	1	
199999	-73.985508	40.768793	1	

```

data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        200000 non-null   int64  
 1   key               200000 non-null   object  
 2   fare_amount       200000 non-null   float64 
 3   pickup_datetime   200000 non-null   object  
 4   pickup_longitude  200000 non-null   float64 
 5   pickup_latitude   200000 non-null   float64 
 6   dropoff_longitude 199999 non-null   float64 
 7   dropoff_latitude  199999 non-null   float64 
 8   passenger_count   200000 non-null   int64  
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB

data.describe()

      Unnamed: 0    fare_amount  pickup_longitude  pickup_latitude
\count  2.000000e+05  200000.000000  200000.000000  200000.000000
mean    2.771250e+07    11.359955     -72.527638    39.935885
std     1.601382e+07    9.901776     11.437787    7.720539
min     1.000000e+00   -52.000000   -1340.648410   -74.015515
25%    1.382535e+07    6.000000    -73.992065    40.734796
50%    2.774550e+07    8.500000    -73.981823    40.752592
75%    4.155530e+07   12.500000    -73.967154    40.767158
max    5.542357e+07   499.000000    57.418457   1644.421482

      dropoff_longitude  dropoff_latitude  passenger_count
count  199999.000000  199999.000000  200000.000000
mean   -72.525292    39.923890    1.684535
std    13.117408     6.794829    1.385997
min   -3356.666300   -881.985513   0.000000
25%   -73.991407    40.733823   1.000000
50%   -73.980093    40.753042   1.000000
75%   -73.963658    40.768001   2.000000
max   1153.572603   872.697628   208.000000

data.shape

```

```
(200000, 9)

data.columns
Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
       'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
       'dropoff_latitude', 'passenger_count'],
      dtype='object')

data = data.drop(['Unnamed: 0', 'key'], axis = 1)

data.shape
(200000, 7)

data

      fare_amount      pickup_datetime  pickup_longitude \
0            7.5  2015-05-07 19:52:06 UTC      -73.999817
1            7.7  2009-07-17 20:04:56 UTC      -73.994355
2           12.9  2009-08-24 21:45:00 UTC      -74.005043
3            5.3  2009-06-26 08:22:21 UTC      -73.976124
4           16.0  2014-08-28 17:47:00 UTC      -73.925023
...
199995         3.0  2012-10-28 10:49:00 UTC      -73.987042
199996         7.5  2014-03-14 01:09:00 UTC      -73.984722
199997        30.9  2009-06-29 00:42:00 UTC      -73.986017
199998        14.5  2015-05-20 14:56:25 UTC      -73.997124
199999        14.1  2010-05-15 04:08:00 UTC      -73.984395

      pickup_latitude  dropoff_longitude  dropoff_latitude
passenger_count
0            40.738354      -73.999512        40.723217
1
1            40.728225      -73.994710        40.750325
1
2            40.740770      -73.962565        40.772647
1
3            40.790844      -73.965316        40.803349
3
4            40.744085      -73.973082        40.761247
5
...
...
199995        40.739367      -73.986525        40.740297
1
199996        40.736837      -74.006672        40.739620
1
199997        40.756487      -73.858957        40.692588
2
199998        40.725452      -73.983215        40.695415
```

```

1
199999      40.720077      -73.985508      40.768793
1

[200000 rows x 7 columns]

data['month'] = data['pickup_datetime']
data

      fare_amount      pickup_datetime  pickup_longitude \
0            7.5  2015-05-07 19:52:06 UTC      -73.999817
1            7.7  2009-07-17 20:04:56 UTC      -73.994355
2           12.9  2009-08-24 21:45:00 UTC      -74.005043
3            5.3  2009-06-26 08:22:21 UTC      -73.976124
4           16.0  2014-08-28 17:47:00 UTC      -73.925023
...
199995         ...  2012-10-28 10:49:00 UTC      -73.987042
199996         ...  2014-03-14 01:09:00 UTC      -73.984722
199997         ...  2009-06-29 00:42:00 UTC      -73.986017
199998         ...  2015-05-20 14:56:25 UTC      -73.997124
199999         ...  2010-05-15 04:08:00 UTC      -73.984395

      pickup_latitude  dropoff_longitude  dropoff_latitude
passenger_count \
0            40.738354      -73.999512      40.723217
1
1            40.728225      -73.994710      40.750325
1
2            40.740770      -73.962565      40.772647
1
3            40.790844      -73.965316      40.803349
3
4            40.744085      -73.973082      40.761247
5
...
199995         ...  40.739367      -73.986525      40.740297
1
199996         ...  40.736837      -74.006672      40.739620
1
199997         ...  40.756487      -73.858957      40.692588
2
199998         ...  40.725452      -73.983215      40.695415
1
199999         ...  40.720077      -73.985508      40.768793
1

          month
0  2015-05-07 19:52:06 UTC
1  2009-07-17 20:04:56 UTC

```

```

2      2009-08-24 21:45:00 UTC
3      2009-06-26 08:22:21 UTC
4      2014-08-28 17:47:00 UTC
...
199995  2012-10-28 10:49:00 UTC
199996  2014-03-14 01:09:00 UTC
199997  2009-06-29 00:42:00 UTC
199998  2015-05-20 14:56:25 UTC
199999  2010-05-15 04:08:00 UTC

[200000 rows x 8 columns]

data['month'] = data['month'].str.slice(start = 5, stop = 7)
data

      fare_amount      pickup_datetime  pickup_longitude \
0          7.5  2015-05-07 19:52:06 UTC           -73.999817
1          7.7  2009-07-17 20:04:56 UTC           -73.994355
2         12.9  2009-08-24 21:45:00 UTC           -74.005043
3          5.3  2009-06-26 08:22:21 UTC           -73.976124
4         16.0  2014-08-28 17:47:00 UTC           -73.925023
...
199995       ...  2012-10-28 10:49:00 UTC           -73.987042
199996       ...  2014-03-14 01:09:00 UTC           -73.984722
199997       ...  2009-06-29 00:42:00 UTC           -73.986017
199998       ...  2015-05-20 14:56:25 UTC           -73.997124
199999       ...  2010-05-15 04:08:00 UTC           -73.984395

      pickup_latitude  dropoff_longitude  dropoff_latitude
passenger_count \
0            40.738354           -73.999512           40.723217
1
1            40.728225           -73.994710           40.750325
1
2            40.740770           -73.962565           40.772647
1
3            40.790844           -73.965316           40.803349
3
4            40.744085           -73.973082           40.761247
5
...
...
199995       ...           -73.986525           40.740297
1
199996       ...           -74.006672           40.739620
1
199997       ...           -73.858957           40.692588
2
199998       ...           -73.983215           40.695415
1

```

```
199999      40.720077      -73.985508      40.768793
1
```

```
      month
0        05
1        07
2        08
3        06
4        08
...
199995     10
199996     03
199997     06
199998     05
199999     05
```

```
[200000 rows x 8 columns]
```

```
data['hour'] = data['pickup_datetime']
data
```

```
      fare_amount      pickup_datetime  pickup_longitude \
0          7.5  2015-05-07 19:52:06 UTC      -73.999817
1          7.7  2009-07-17 20:04:56 UTC      -73.994355
2         12.9  2009-08-24 21:45:00 UTC      -74.005043
3          5.3  2009-06-26 08:22:21 UTC      -73.976124
4         16.0  2014-08-28 17:47:00 UTC      -73.925023
...
199995       3.0  2012-10-28 10:49:00 UTC      -73.987042
199996       7.5  2014-03-14 01:09:00 UTC      -73.984722
199997      30.9  2009-06-29 00:42:00 UTC      -73.986017
199998      14.5  2015-05-20 14:56:25 UTC      -73.997124
199999      14.1  2010-05-15 04:08:00 UTC      -73.984395
```

```
      pickup_latitude  dropoff_longitude  dropoff_latitude
passenger_count \
0            40.738354      -73.999512      40.723217
1
1            40.728225      -73.994710      40.750325
1
2            40.740770      -73.962565      40.772647
1
3            40.790844      -73.965316      40.803349
3
4            40.744085      -73.973082      40.761247
5
...
199995       40.739367      -73.986525      40.740297
1
```

```
199996      40.736837     -74.006672      40.739620
1
199997      40.756487     -73.858957      40.692588
2
199998      40.725452     -73.983215      40.695415
1
199999      40.720077     -73.985508      40.768793
1
```

```
    month          hour
0      05 2015-05-07 19:52:06 UTC
1      07 2009-07-17 20:04:56 UTC
2      08 2009-08-24 21:45:00 UTC
3      06 2009-06-26 08:22:21 UTC
4      08 2014-08-28 17:47:00 UTC
...
199995     10 2012-10-28 10:49:00 UTC
199996     03 2014-03-14 01:09:00 UTC
199997     06 2009-06-29 00:42:00 UTC
199998     05 2015-05-20 14:56:25 UTC
199999     05 2010-05-15 04:08:00 UTC
```

[200000 rows x 9 columns]

```
data['hour'] = data['hour'].str.slice(start = 11, stop = 13)
data
```

```
    fare_amount      pickup_datetime  pickup_longitude \
0        7.5 2015-05-07 19:52:06 UTC      -73.999817
1        7.7 2009-07-17 20:04:56 UTC      -73.994355
2       12.9 2009-08-24 21:45:00 UTC      -74.005043
3        5.3 2009-06-26 08:22:21 UTC      -73.976124
4       16.0 2014-08-28 17:47:00 UTC      -73.925023
...
199995       3.0 2012-10-28 10:49:00 UTC      -73.987042
199996       7.5 2014-03-14 01:09:00 UTC      -73.984722
199997      30.9 2009-06-29 00:42:00 UTC      -73.986017
199998      14.5 2015-05-20 14:56:25 UTC      -73.997124
199999      14.1 2010-05-15 04:08:00 UTC      -73.984395
```

```
    pickup_latitude  dropoff_longitude  dropoff_latitude
passenger_count \
0            40.738354      -73.999512      40.723217
1
1            40.728225      -73.994710      40.750325
1
2            40.740770      -73.962565      40.772647
1
3            40.790844      -73.965316      40.803349
3
```

```
4          40.744085        -73.973082        40.761247
5
...
199995      40.739367        -73.986525        40.740297
1
199996      40.736837        -74.006672        40.739620
1
199997      40.756487        -73.858957        40.692588
2
199998      40.725452        -73.983215        40.695415
1
199999      40.720077        -73.985508        40.768793
1
```

```
    month hour
0      05   19
1      07   20
2      08   21
3      06   08
4      08   17
...
199995    10   10
199996    03   01
199997    06   00
199998    05   14
199999    05   04
```

```
[200000 rows x 9 columns]
```

```
data = data.drop(['pickup_datetime'], axis = 1)
data
```

```
    fare_amount  pickup_longitude  pickup_latitude
dropoff_longitude \
0            7.5        -73.999817        40.738354
73.999512
1            7.7        -73.994355        40.728225
73.994710
2           12.9        -74.005043        40.740770
73.962565
3            5.3        -73.976124        40.790844
73.965316
4           16.0        -73.925023        40.744085
73.973082
...
199995      3.0        -73.987042        40.739367
73.986525
199996      7.5        -73.984722        40.736837
```

```

74.006672
199997      30.9       -73.986017      40.756487      -
73.858957
199998      14.5       -73.997124      40.725452      -
73.983215
199999      14.1       -73.984395      40.720077      -
73.985508

```

	dropoff_latitude	passenger_count	month	hour
0	40.723217		1	05 19
1	40.750325		1	07 20
2	40.772647		1	08 21
3	40.803349		3	06 08
4	40.761247		5	08 17
..
199995	40.740297		1	10 10
199996	40.739620		1	03 01
199997	40.692588		2	06 00
199998	40.695415		1	05 14
199999	40.768793		1	05 04

[200000 rows x 8 columns]

data.describe()

	fare_amount	pickup_longitude	pickup_latitude
dropoff_longitude \ count	200000.000000	200000.000000	200000.000000
199999.000000			
mean	11.359955	-72.527638	39.935885
72.525292			
std	9.901776	11.437787	7.720539
13.117408			
min	-52.000000	-1340.648410	-74.015515
3356.666300			
25%	6.000000	-73.992065	40.734796
73.991407			
50%	8.500000	-73.981823	40.752592
73.980093			
75%	12.500000	-73.967154	40.767158
73.963658			
max	499.000000	57.418457	1644.421482
1153.572603			

	dropoff_latitude	passenger_count
count	199999.000000	200000.000000
mean	39.923890	1.684535
std	6.794829	1.385997
min	-881.985513	0.000000
25%	40.733823	1.000000

```

50%          40.753042      1.000000
75%          40.768001      2.000000
max         872.697628    208.000000

data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 8 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   fare_amount       200000 non-null   float64
 1   pickup_longitude  200000 non-null   float64
 2   pickup_latitude   200000 non-null   float64
 3   dropoff_longitude 199999 non-null   float64
 4   dropoff_latitude  199999 non-null   float64
 5   passenger_count   200000 non-null   int64  
 6   month              200000 non-null   object 
 7   hour               200000 non-null   object 
dtypes: float64(5), int64(1), object(2)
memory usage: 12.2+ MB

data["dropoff_longitude"]=
data["dropoff_longitude"].fillna(data['dropoff_longitude'].mean())
data["dropoff_latitude"]=
data["dropoff_latitude"].fillna(data['dropoff_latitude'].mean())

# data.dropna(inplace=True)----- Drop complete row

data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 8 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   fare_amount       200000 non-null   float64
 1   pickup_longitude  200000 non-null   float64
 2   pickup_latitude   200000 non-null   float64
 3   dropoff_longitude 200000 non-null   float64
 4   dropoff_latitude  200000 non-null   float64
 5   passenger_count   200000 non-null   int64  
 6   month              200000 non-null   object 
 7   hour               200000 non-null   object 
dtypes: float64(5), int64(1), object(2)
memory usage: 12.2+ MB

def haversine (lon_1, lon_2, lat_1, lat_2):
    lon_1, lon_2, lat_1, lat_2 = map(np.radians, [lon_1, lon_2, lat_1,
lat_2])
    diff_lon = lon_2 - lon_1

```

```

    diff_lat = lat_2 - lat_1
    km = 2 * 6371 * np.arcsin(np.sqrt(np.sin(diff_lat/2.0)**2 +
np.cos(lat_1) * np.cos(lat_2) * np.sin(diff_lon/2.0)**2))
    return km

data['distance'] =
haversine(data['pickup_longitude'],data['dropoff_longitude'],data['pic
kup_latitude'],data['dropoff_latitude'])

data.describe()

      fare_amount  pickup_longitude  pickup_latitude
dropoff_longitude \
count  200000.000000           200000.000000           200000.000000
200000.000000
mean       11.359955          -72.527638          39.935885
72.525292
std        9.901776          11.437787          7.720539
13.117375
min      -52.000000         -1340.648410         -74.015515
3356.666300
25%       6.000000          -73.992065          40.734796
73.991407
50%       8.500000          -73.981823          40.752592
73.980093
75%       12.500000         -73.967154          40.767158
73.963658
max      499.000000          57.418457          1644.421482
1153.572603

      dropoff_latitude  passenger_count  distance
count  200000.000000           200000.000000  200000.000000
mean       39.923890          1.684535          20.856014
std        6.794812          1.385997          382.963800
min      -881.985513          0.000000          0.000000
25%       40.733823          1.000000          1.215222
50%       40.753042          1.000000          2.121005
75%       40.768001          2.000000          3.875248
max      872.697628          208.000000          16409.239135

data.replace(to_replace = 0, value = data['passenger_count'].mean(),
inplace=True)
data.replace(to_replace = 0, value = data['distance'].mean(),
inplace=True)
data[data['fare_amount'] <= 0] = data['fare_amount'].mean()

data.describe()

      fare_amount  pickup_longitude  pickup_latitude
dropoff_longitude \
count  200000.000000           200000.000000           200000.000000

```

```
200000.000000
mean      11.362407      -72.488497      39.965242      -
72.486713
std       9.896653      11.666136      7.562049
13.314895
min      0.010000     -1340.648410     -74.015515      -
3356.666300
25%      6.000000      -73.992065      40.734785      -
73.991407
50%      8.500000      -73.981821      40.752590      -
73.980091
75%      12.500000      -73.967148      40.767157      -
73.963653
max      499.000000      57.418457      1644.421482
1153.572603
```

```
dropoff_latitude  passenger_count  distance
count      200000.000000      200000.000000      200000.000000
mean      39.953240      1.691287      20.860885
std       6.614871      1.385145      382.473053
min      -881.985513      1.000000      0.000084
25%      40.733817      1.000000      1.305244
50%      40.753040      1.000000      2.121431
75%      40.767999      2.000000      3.876122
max      872.697628      208.000000      16409.239135
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 200000 entries, 0 to 199999
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	fare_amount	200000 non-null	float64
1	pickup_longitude	200000 non-null	float64
2	pickup_latitude	200000 non-null	float64
3	dropoff_longitude	200000 non-null	float64
4	dropoff_latitude	200000 non-null	float64
5	passenger_count	200000 non-null	float64
6	month	200000 non-null	object
7	hour	200000 non-null	object
8	distance	200000 non-null	float64

```
dtypes: float64(7), object(2)
```

```
memory usage: 13.7+ MB
```

```
data
```

```
fare_amount  pickup_longitude  pickup_latitude
dropoff_longitude \
0           7.5            -73.999817          40.738354      -
```

```

73.999512
1          7.7      -73.994355      40.728225      -
73.994710
2         12.9      -74.005043      40.740770      -
73.962565
3          5.3      -73.976124      40.790844      -
73.965316
4         16.0      -73.925023      40.744085      -
73.973082
...
...
199995      3.0      -73.987042      40.739367      -
73.986525
199996      7.5      -73.984722      40.736837      -
74.006672
199997     30.9      -73.986017      40.756487      -
73.858957
199998     14.5      -73.997124      40.725452      -
73.983215
199999     14.1      -73.984395      40.720077      -
73.985508


      dropoff_latitude  passenger_count  month  hour  distance
0           40.723217            1.0    05   19   1.683323
1           40.750325            1.0    07   20   2.457590
2           40.772647            1.0    08   21   5.036377
3           40.803349            3.0    06   08   1.661683
4           40.761247            5.0    08   17   4.475450
...
...
199995      40.740297            1.0    10   10   0.112210
199996      40.739620            1.0    03   01   1.875050
199997      40.692588            2.0    06   00  12.850319
199998      40.695415            1.0    05   14   3.539715
199999      40.768793            1.0    05   04   5.417783

[200000 rows x 9 columns]

data.plot(kind = "box", subplots = True, layout =
(6,2), figsize=(15,20))

fare_amount          AxesSubplot(0.125,0.772143;0.352273x0.107857)
pickup_longitude    AxesSubplot(0.547727,0.772143;0.352273x0.107857)
pickup_latitude      AxesSubplot(0.125,0.642714;0.352273x0.107857)
dropoff_longitude   AxesSubplot(0.547727,0.642714;0.352273x0.107857)
dropoff_latitude     AxesSubplot(0.125,0.513286;0.352273x0.107857)
passenger_count      AxesSubplot(0.547727,0.513286;0.352273x0.107857)
distance             AxesSubplot(0.125,0.383857;0.352273x0.107857)
dtype: object

```



```
# Using the InterQuartile Range to fill the values
def remove_outlier(df1 , col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_whisker = Q1 - 1.5 * IQR
    upper_whisker = Q3 + 1.5 * IQR
    data[col] = np.clip(df1[col] , lower_whisker , upper_whisker)
    return df1


def treat_outliers_all(df1 , col_list):
    for c in col_list:
        df1 = remove_outlier(data , c)
    return df1

cols = ['fare_amount', 'pickup_longitude', 'pickup_latitude',
        'dropoff_longitude', 'dropoff_latitude', 'passenger_count',
```

```
'distance']
data = treat_outliers_all(data , cols)

data.plot(kind = "box",subplots = True, layout =
(6,2),figsize=(15,20))

fare_amount          AxesSubplot(0.125,0.772143;0.352273x0.107857)
pickup_longitude    AxesSubplot(0.547727,0.772143;0.352273x0.107857)
pickup_latitude      AxesSubplot(0.125,0.642714;0.352273x0.107857)
dropoff_longitude   AxesSubplot(0.547727,0.642714;0.352273x0.107857)
dropoff_latitude     AxesSubplot(0.125,0.513286;0.352273x0.107857)
passenger_count      AxesSubplot(0.547727,0.513286;0.352273x0.107857)
distance            AxesSubplot(0.125,0.383857;0.352273x0.107857)
dtype: object
```



```
data.shape
```

```
(200000, 9)
```

```
data.isnull().sum()


fare_amount      0
pickup_longitude 0
pickup_latitude   0
dropoff_longitude 0
dropoff_latitude  0
passenger_count   0
month            0
hour             0
distance          0
dtype: int64
```

VISUALIZATION

```
import seaborn as sns
import matplotlib.pyplot as plt


plt.scatter(data['pickup_longitude'], data['fare_amount'])
plt.xlabel("pickup_longitude")
plt.ylabel("fare_amount")

Text(0, 0.5, 'fare_amount')
```




```
plt.scatter(data['pickup_latitude'], data['fare_amount'])
plt.xlabel("pickup_latitude")
plt.ylabel("fare_amount")
```

```
Text(0, 0.5, 'fare_amount')
```




```
plt.scatter(data['dropoff_longitude'], data['fare_amount'])  
plt.xlabel("dropoff_longitude")  
plt.ylabel("fare_amount")
```

```
Text(0, 0.5, 'fare_amount')
```




```
plt.scatter(data['dropoff_longitude'], data['fare_amount'])
plt.xlabel("dropoff_longitude")
plt.ylabel("fare_amount")
Text(0, 0.5, 'fare_amount')
```



```

plt.scatter(data['distance'], data['fare_amount'])
plt.xlabel("distance")
plt.ylabel("fare_amount")
Text(0, 0.5, 'fare_amount')

```



```

corr_matrix = data.corr()
corr_matrix

            fare_amount  pickup_longitude  pickup_latitude \
fare_amount      1.000000      0.154189     -0.110927
pickup_longitude    0.154189      1.000000      0.258704
pickup_latitude     -0.110927      0.258704      1.000000
dropoff_longitude    0.218764      0.425930      0.048264
dropoff_latitude     -0.125988      0.072689      0.515985
passenger_count       0.014649     -0.012819     -0.013306
distance              0.858832      0.134118     -0.082360

            dropoff_longitude  dropoff_latitude
passenger_count \
fare_amount           0.218764      -0.125988
0.014649
pickup_longitude        0.425930      0.072689
0.012819
pickup_latitude         0.048264      0.515985
0.013306
dropoff_longitude        1.000000      0.244958
0.008774

```


```

dropoff_latitude          0.244958      1.000000      -
0.006609
passenger_count           -0.008774     -0.006609      1.000000
distance                  0.224214     -0.072548      0.009658

                           distance
fare_amount               0.858832
pickup_longitude          0.134118
pickup_latitude            -0.082360
dropoff_longitude          0.224214
dropoff_latitude           -0.072548
passenger_count            0.009658
distance                  1.000000

sns.heatmap(corr_matrix, annot = True)
<matplotlib.axes._subplots.AxesSubplot at 0x23d90029cd0>

```



Splitting the dataset

```

X = data.iloc[:, 1:]
y = data.iloc[:, 0]
print(X.shape, y.shape)

(200000, 8) (200000,)

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.15, random_state=42)
print(X_train.shape, y_train.shape, X_test.shape, y_test.shape)

(170000, 8) (170000,) (30000, 8) (30000,)

```

Linear Regression

```

from sklearn.linear_model import LinearRegression
linear_regression = LinearRegression().fit(X_train, y_train)

linear_regression.score(X_test, y_test)

0.7316136867718988

y_pred = linear_regression.predict(X_test)

result = pd.DataFrame()
result['Actual'], result['Predicted'] = y_test, y_pred
result.sample(10)

      Actual   Predicted
63108    11.0    9.709085
82890     4.0    4.895952
102306    12.1    9.741012
92646    14.0   13.664903
47424    12.9    9.483992
172670     7.5    7.203042
133321    14.0   13.618706
177665     8.1    9.410597
58919     13.3   16.056232
64404      6.9    7.110601

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test,
y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test,
y_pred))
print('Root Mean Squared Error:',
np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('R Squared (R2):', np.sqrt(metrics.r2_score(y_test, y_pred)))

Mean Absolute Error: 1.8763278733128748
Mean Squared Error: 7.9583356668996945
Root Mean Squared Error: 2.821052226900398
R Squared (R2): 0.8553441919905101

```

Random Forest

```
from sklearn.ensemble import RandomForestRegressor
random_forest = RandomForestRegressor(n_estimators = 10, random_state = 42)

random_forest.fit(X_train, y_train)
RandomForestRegressor(n_estimators=10, random_state=42)
random_forest.score(X_test, y_test)
0.7525809054770123

y_pred = random_forest.predict(X_test)

result = pd.DataFrame()
result['Actual'], result['Predicted'] = y_test, y_pred
result.sample(10)

      Actual   Predicted
108389    7.70      9.50
16131     9.00      9.33
138295   22.25     19.76
32169    11.50      9.28
172860     5.00      4.77
20218    11.70     14.53
136756     7.70     10.04
148663     5.30      5.20
13249     2.90      7.17
78967    11.00      9.70

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test,
y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test,
y_pred))
print('Root Mean Squared Error:',
np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
print('R Squared (R2):', np.sqrt(metrics.r2_score(y_test, y_pred)))

Mean Absolute Error: 1.7530684681716462
Mean Squared Error: 7.336604392865712
Root Mean Squared Error: 2.7086166936031595
R Squared (R2): 0.8675142105331833
```